

# Background on Computation/AI

## Agent acting intelligently in its own environment (†)

- Actions appropriate for goals
- Flexible to changing environments & goals
- Learns from Experience
- Appropriate choices for limitations, finite computation

## Symbol-System Hypothesis

A physical symbol system has the necessary and sufficient means for general intelligent action

- Necessity: Anything capable of intelligent action is a physical symbol system
- Sufficiency: Any (sufficiently sophisticated) PSS is capable of intelligent action
- Reasoning is Symbol Manipulation
- Symbols are 1 and 0 of computers, in which case it only means intelligence can be digitized

## Doubts

- The brain is not merely a computer, *computation* is not a complete model for intelligence

## Church Turing Thesis and its relevance to AI

- Any symbol manipulation possible on a Turing Machine
- Combined with *SSH*, any Intelligent Action possible on a Turing Machine
- Computation turns into Graph Search (express problem as a graph)

## Turing machine

A theoretical machine to reason about computation. Reads and writes symbols to and from a tape (†'s environment). A 6-tuple of:

- $\Sigma$  : Alphabet of Symbols
- $Q$  : Set of possible internal states
- $Q_0 \in Q$  : Initial State
- $\epsilon \in \Sigma$  : Blank symbol
- $A$  : Accepting/Final States

- $\delta \subseteq (Q \setminus A \times \Sigma) \times (Q \times \Sigma \times \{L, R\})$  : Relation on State-Symbol pairs, mapping to State-Symbol-Left/Right Movement

Example:  $(q, \sigma), (r, \alpha, L)$  would encode, if in state  $q$  and a  $\sigma$  is read, move to state  $r$ , write an  $\alpha$  to the tape, and move Left.

- **Deterministic Turing Machine (DTM)** has one State-Symbol-Left/Right triple per State-Symbol pair (it's a function)
- **Non-Deterministic Turing Machine (NTM)** can have more than one S-S-LR per S-S pair.
- †: The environment is the tape, and the agent is the TM. If the TM reaches a state in  $A$ , then the action of the agent has been complete/solution to given problem found.

## Non-determinism

- Turing machine defines next state depending on current state, read symbol. An NTM has more than one S-S-LR triple per S-S pair.
- With computation as graph search
 

```
search(Node) :- goal(Node).
search(Node) :- arc(Node, Next), search(Next).
```

there may exist more than one *Next* for some *Node*.
- **Don't know (Prolog does this)** : If one choice doesn't lead to solution another might
  - Choose
- **Don't care (Paralog)** : If one selection doesn't lead to a solution, no point in trying others
  - select

†: Problem may be a non-deterministic one, and so needs an NTM to implement the Intelligent Agent to solve it feasibly.

## $P$ vs $NP$

### Cobham's Thesis

Feasible computation is defined as being solved by a Deterministic Turing Machine in Polynomial time (it's in  $P$ ).

$P := \{\text{problems solved by a Deterministic Turing Machine in Polynomial Time}\}$

$NP := \{\text{problems solved by a Non-deterministic Turing Machine in Polynomial Time}\}$

$P$  is clearly  $\subseteq NP$  as all DTMS can just be an NTM with only one triple in its  $\delta$  relation (in which  $\delta$  then defines a function).

## SAT

For some boolean expression  $\phi$ , of variables  $x_1, x_2, \dots, x_n$ , find an assignment makes  $\phi$  evaluate to True. Checking that it's True in  $P$  is easy, linear in  $n$ . Finding is hard due to non-determinism, many assignment to check.

- **Cook-Levin Theorem** says  $SAT \in P \iff P = NP$ .
- **CSAT** :  $\phi$  is a *conjunction* of *clauses* where a *clause* is a disjunction of *literals* and a literal is either a non-negated variable  $x_i$  (positive) or a negated variable  $\neg x_i$  (negative).
- **k-SAT** : Says each clause has  $k$  literals
- **3-SAT** : Says each clause has 3 literals. Is as hard as SAT, but 2-SAT is in  $P$ .
  - **horn-SAT** : A conjunction of *horn* clause s, where a horn clause has at most 1 positive literal. Linear.

## Halting Problem

Given a program  $P$ , and data  $D$  return 1 if  $P$  halts on  $D$ , otherwise 0 (if it loops indefinitely). It is undecidable.

### Proof (by contadiction)

- Assume for contadiction's sake  $\exists$  program  $halt(P, D)$  that returns 1 iff  $P$  halts on  $D$ , otherwise 0.
- Now construct a new program/string  $Z$

```
def Z(String x)
  if halt(x, x) then
    loop forever
  else
    halt
  end
end
```

and run it on itself, i.e.  $Z(Z)$ . There are 2 cases:

1.  $Z$  halts on  $Z$ . Then the call to `Halt( $Z$ ,  $Z$ )` will return `True`, so  $Z$  loops forever on  $Z$ .
    - Contradiction
  2.  $Z$  loops forever on  $Z$ . Then the call to `Halt( $Z$ ,  $Z$ )` returns `False`, so  $Z$  halts on  $Z$ 
    - Contradiction.
- Conclude that *halt* cannot exist, so there is no general method to decide if some  $P$  will halt on some  $D$ .

Prolog consequence is that there's no general algorithm to detect loops caused by KBs such as

```
p :- q.
q :- p.
```

```
a :- a.
```

etc.

## Church-Turing Thesis

A function is effectively calculable if its values can be found by some purely mechanical process.

### Halting Problem implications:

Can define functions which are not computable.

- `busy_beaver( $n$ )`: given a TM with  $n$  possible states, how many symbols can it write before halting when run with no input?

Can't get an upper bound on this without solving the halting problem. Since *CTT* says you can compute anything on a *TM* this is uncomputable by any method.

## Cantor's Theorem

$$\forall \omega, |2^\omega| > |\omega|$$

TODO: Prove this

# Knowledge Representation and Reasoning

## Representation and Reasoning System

### Definitions

- **Formal Language** : legal sentences
- **Semantics** : meaning of the symbols
- **Reasoning theory/proof procedure** *nondeterministic* specification for how to produce and answer

### Implementation

- **Language Parser** : Sentences  $\rightarrow$  Data Structures
- **Reasoning Procedure** : implementation of reasoning theory, search strategy
  - Does *not* reflect semantics (It's a symbol-system manipulation)

### Datalog

**propositional definite clause** : one of these?

- variable : starts with upper case
- constant : starts with lower case, or is a numeral
- predicate symbol : starts with lower case
- term : variable or a constant
- **atomic symbol (atom)** :  $p$  or  $p(t_1, \dots, t_n)$ , with  $p$  a predicate and each  $t_i$  is a term.
- **definite clause** :  $a \leftarrow b_1 \wedge \dots \wedge b_m$
- **query** :  $?b_1 \wedge \dots \wedge b_m$
- **knowledge base** : set of definite clauses

### Semantics

Meaning of sentences in a language

### Interpretation

What is in the world, symbol-to-real-things-and-relations correspondence

Triple  $I = \langle D, \phi, \pi \rangle$

- **D** : Nonempty domain set. Elements are *individuals*

- $\phi$  : mapping each constant to an individual. A constant  $c$  denotes an individual  $\phi(c)$
- $\pi$  : maps each to  $n$ -ary predicate symbol a relation
  - ie  $\pi : D^n \mapsto \{\text{True}, \text{False}\}$

### Notes

- $D$  can be actual real things, not confined to being in a computer.
- $\pi(p)$  specifies truthiness for the predicate symbol  $p$ , for each  $n$ -tuple of individuals
- if  $p$  has no arguments, then  $\pi(p)$  either True or False.

### Truth in Interpretation

- $c$  denotes in  $I$  the individual  $\phi(c)$ .
- **Ground** (variable free) atom  $p(t_1, \dots, t_n)$ , in **Interpretation**  $I$  is
  - **True** if  $\pi(p)(t'_1, \dots, t'_n) = \text{True}$
  - **False** if  $\pi(p)(t'_1, \dots, t'_n) = \text{False}$ 
    - \* Where  $t_i$  denotes  $t'_i$  in interpretation  $I$
- Ground clause  $h \leftarrow b_1 \wedge \dots \wedge b_m$  is **False in**  $I$  if  $h$  is False in  $I$  and each  $b_i$  is True in  $I$ .
  - Otherwise it's *True* in  $I$ .

### Models, Logical Consequence

- A Knowledge Base  $KB$  is True in interpretation  $I$  iff every clause in  $KB$  is True in interpretation  $I$  iff every clause in  $KB$  is True in  $I$ .
- a **model** is an interpretation in which all clauses are True
- $g$  is a **logical consequence** of  $KB$  ( $KB \models g$ ), if  $g$  is True in every models of  $KB$ .
  - ( $KB \models g$ ) if there's no  $I$  such that  $KB$  is True  $\wedge$   $g$  is False.

### For Users

1. Come up with an **intended interpretation**  $I$ 
  - the problem domain
2. Pick constants for the relevant individuals
  - e.g. **shibe** for your pet
3. Pick a predicate symbol for the relations
  - **is\_dog** to denote a constant's individual being a dog
4. Tell it things that are True in  $I$ 
  - build up the knowledge base by **axiomatizing the domain** )
  - **is\_dog(X) :- barks(X).**
  - **barks(shibe).**, etc.

5. Ask it things
  - `? is_dog(shibe) → yes.`
6. Now if  $KB \models g$ , then  $g$  must be True in  $I$ 
  - Your pet must indeed be a dog.

### For Computers

- Knows nothing about the interpretation, only  $KB$ .
  - What's a dog? What is barks?
- *Can* determine if some  $g$  is an LC of  $KB$ , so if  $KB \models g$ , then it's True in  $I$ .
- If  $\neg(KB \models g)$ , then  $\exists$  some Interpretation in which  $g$  is False. This could be the intended interpretation  $I$ .

## Proofs

## Search

Basic graph searching (Depth first)

```
arc(Node, Next) :- % Something that relates Node to Next.
search(Node) :- goal(Node).
search(Node) :- arc(Node, Next), search(Next).
```

- Nondeterminism if `arc/2` has multiple solutions
  - Choose the best one (A\*, Best first, etc)
- Computation eliminates the non-determinism
- Can bound the number of calls to `arc`, the number of search iterations.

## Frontier Search

- With a **graph**, **start nodes** and **goal nodes**, incrementally explore paths from start nodes, hoping to reach goal nodes.
- Maintain a **frontier** of paths from start that have been explored.
- Search is complete once frontier hits a goal.
- How the frontier expands (how the child nodes of the current frontier are inserted to the frontier) can vary, defines the **search strategy**

```
frontier := { s : s is a start node }
until frontier.empty
  select and remove path <n_0, ... ,n_k> from frontier;
  if goal(n_k)
    return <n_0, ... , n_k>;
  end
```

```

    for each neighbor n of n_k;
        add <n_0, ... , n_k, n> to frontier
    end
end

```

## Prolog Examples

```

search(Start) :- frontier_search([Start]).

frontier_search([Node|_]) :- goal(Node).
frontier_search([Node|Rest]) :-
    findall(Next, arc(Node, Next), Children),
    add_to_frontier(Children, Rest, New), % add the neighbours to the frontier
    frontier_search(New).

```

The strategy depends on how `add_to_frontier/3` is defined:

Depth-first, empty the `Children` before adding the `New`:

```

add_to_frontier([], Rest, Rest).
add_to_frontier([H|T], Rest, [H, New]) :- add_to_frontier(T, Rest, New)

```

Breadth-first, empty the `New` before adding the `Children`:

```

add_to_frontier(Children, [], Children).
add_to_frontier(Children, [H|T], Children) add_to_frontier(Children, T, New).

```

Bounded Depth First by adding an iteration-counting term:

```

bounded_search(Node, _) :- goal(Node).
bounded_search(Node, s(B)) :- bounded_search(B).

```

Iterative deepening is depth first with a variable bound:

```

iterative_deepening(Node) :- bound(B), bs(Next, B).

```

```

bound(0).
bound(s(B)) :- bound(B). % easily adjustable

```

## Knowledge Representation and Reasoning

### Representation and Reasoning System

#### Definitions

- **Formal Language** : legal sentences



- **Semantics** : meaning of the symbols
- **Reasoning theory/proof procedure** *nondeterministic* specification for how to produce and answer

## Implementation

- **Language Parser** : Sentences  $\rightarrow$  Data Structures
- **Reasoning Procedure** : implementation of reasoning theory, search strategy
  - Does *not* reflect semantics (It's a symbol-system manipulation)

## Datalog

**propositional definite clause** : one of these?

- variable : starts with upper case
- constant : starts with lower case, or is a numeral
- predicate symbol : starts with lower case
- term : variable or a constant
- **atomic symbol (atom)** :  $p$  or  $p(t_1, \dots, t_n)$ , with  $p$  a predicate and each  $t_i$  is a term.
- **definite clause** :  $a \leftarrow b_1 \wedge \dots \wedge b_m$
- **query** :  $?b_1 \wedge \dots \wedge b_m$
- **knowledge base** : set of definite clauses

## Semantics

Meaning of sentences in a language

## Interpretation

What is in the world, symbol-to-real-things-and-relations correspondence

Triple  $I = \langle D, \phi, \pi \rangle$

- **D** : Nonempty domain set. Elements are *individuals*
- $\phi$  : mapping each constant to an individual. A constant  $c$  denotes an individual  $\phi(c)$
- $\pi$  : maps each to  $n$ -ary predicate symbol a relation
  - ie  $\pi : D^n \mapsto \{\text{True}, \text{False}\}$

## Notes

- $D$  can be actual real things, not confined to being in a computer.
- $\pi(p)$  specifies truthiness for the predicate symbol  $p$ , for each  $n$ -tuple of individuals
- if  $p$  has no arguments, then  $\pi(p)$  either True or False.

## Truth in Interpretation

- $c$  denotes in  $I$  the individual  $\phi(c)$ .
- **Ground** (variable free) atom  $p(t_1, \dots, t_n)$ , in **Interpretation**  $I$  is
  - **True** if  $\pi(p)(t'_1, \dots, t'_n) = \text{True}$
  - **False** if  $\pi(p)(t'_1, \dots, t'_n) = \text{False}$ 
    - \* Where  $t'_i$  denotes  $t_i$  in interpretation  $I$
- Ground clause  $h \leftarrow b_1 \wedge \dots \wedge b_m$  is **False in**  $I$  if  $h$  is False in  $I$  and each  $b_i$  is True in  $I$ .
  - Otherwise it's *True* in  $I$ .

## Models, Logical Consequence

- A Knowledge Base  $KB$  is True in interpretation  $I$  iff every clause in  $KB$  is True in interpretation  $I$  iff every clause in  $KB$  is True in  $I$ .
- a **model** is an interpretation in which all clauses are True
- $g$  is a **logical consequence** of  $KB$  ( $KB \models g$ ), if  $g$  is True in every models of  $KB$ .
  - ( $KB \models g$ ) if there's no  $I$  such that  $KB$  is True  $\wedge$   $g$  is False.

## For Users

1. Come up with an **intended interpretation**  $I$ 
  - the problem domain
2. Pick constants for the relevant individuals
  - e.g. **shibe** for your pet
3. Pick a predicate symbol for the relations
  - **is\_dog** to denote a constant's individual being a dog
4. Tell it things that are True in  $I$ 
  - build up the knowledge base by **axiomatizing the domain** )
  - **is\_dog(X) :- barks(X).**
  - **barks(shibe).**, etc.
5. Ask it things
  - **? is\_dog(shibe) → yes.**
6. Now if  $KB \models g$ , then  $g$  must be True in  $I$ 
  - Your pet must indeed be a dog.

### For Computers

- Knows nothing about the interpretation, only  $KB$ .
  - What's a dog? What is barks?
- *Can* determine if some  $g$  is an LC of  $KB$ , so if  $KB \models g$ , then it's True in  $I$ .
- If  $\neg(KB \models g)$ , then  $\exists$  some Interpretation in which  $g$  is False. This could be the intended interpretation  $I$ .

### Proofs

- **Proof** : mechanical derivation that formula *follows* from KB.
  - $KB \vdash g$  :  $g$  can be derived from  $KB$ .
  - $KB \models g$  :  $g$  is true in *all models* of  $KB$ .
- **soundness** :  $KB \vdash g \Rightarrow KB \models g$
- **completeness** :  $KB \models g \Rightarrow KB \vdash g$

### Bottom Up

A rule of derivation, modus ponens *generalisé*.

- Given  $h \leftarrow b_1 \wedge \dots \wedge b_m$  in the KB, if each  $b_i$  has been derived, then  $h$  can be derived.
- **Forward chaining** on the clause. Also covers  $m = 0$ .

### Procedure

If  $g \in C$  at the end, then  $KB \vdash g$ .

```
C := {}
until (no more clauses can be selected) {
  select clause h <- b_1 && ... && b_n such that
    b_i in C for all i
    and h not in C ;
  C := C union {h};
}
```

### Proof of Soundness:

$KB \vdash g \Rightarrow KB \models g$

- Suppose  $\exists g : KB \vdash g \wedge \neg(KB \models g)$ .
- Let  $h$  be the first atom added to  $C$ , which is not true in every model of  $KB$ .
- Suppose  $h$  is not true in model  $I$  of  $KB$ .
- There must be some clause  $h \leftarrow b_1 \wedge \dots \wedge b_m$

- Now each  $b_i$  is True in  $I$ .  $h$  is False in  $I$ , so the clause is False in  $I$ .
- So  $I$  can't be a model of  $KB$ .
- This is a contradiction, so no such  $g$  exists.

### Fixed point

$C$ , at the end, is a **fixed point**.

Now if we let  $I$  be the interpretation such that every element of the *fixed point* is True, and every other atom is False, then  $I$  is a model of  $KB$ .

- Suppose  $h \leftarrow b_1 \wedge \dots \wedge b_m \in KB$  is False in  $I$ . Then  $h$  is False, and each  $b_i$  is true in  $I$ . So we can, by the method, add  $h$  to  $C$ .
- This contradicts  $C$  being a fixed point.
- The  $I$  is a **minimal model**.

### Proof of Completeness

- $KB \models g \Rightarrow KB \vdash g$ .
- Suppose  $KB \models g$ . Then  $g$  is true in all models of  $KB$ . Thus  $g$  is true in the minimal model.
- Thus  $g$  is generated by the bottom up algorithm.
- Thus  $KB \vdash g$ .

### Top Down

- Go back from a query to see if it's a logical consequence of the KB.
- **Answer Clause**  $\alpha$  is  $\text{Yes} \leftarrow a_1 \wedge \dots \wedge c_m$
- **SLD Resolution** of  $\alpha$  with atom  $a_i$  is  $\alpha$  with  $a_i$  substituted for the clauses of  $a_i$
- **Answer** is an answer clause with  $m = 0$ , i.e.  $\text{Yes} \leftarrow$ .
- **derivation** is a sequence of answer clauses,  $\gamma_0, \gamma_1, \dots, \gamma_n$ 
  - Each  $\gamma_i$  is the resolution of  $\gamma_{i-1}$  with some clause in the KB.
  - $\gamma_n$  is the 'final' answer.

### Procedure

To solve  $ac$ .

```
ac := "yes <- q_1 && ... && q_k"
until ac is an answer (i.e. until ac matches "yes <-")
  select conjunct a_i from the body of a_c;
  choose C from the KB that has a_i at its head;
```

```
    replace a_i in body of ac with body of C;
end
```

There's nondeterminism in the choice of C here:

- **Don't care** : If one doesn't lead to solution, none of the others will.
- **Don't know** : If one choice doesn't lead to the solution, others might.

## Knowledge Representation

- How to represent "*Coco is a Shiba Inu*"
- Could do something like `shibe(coco)`
  - "Who are the shibes?" easy to solve
- Or `breed(coco, shibe)`
  - "What breed is Coco"
  - "What dogs are the Shiba Inu"
  - ~"What property is"Shiba Inu;~~ It's a **breed**, but we can't resolve this easily.
- Solution: `prop(coco, breed, shibe)`. Now we can solve it all with the usual strategies. Called **object-attribute-value** representation.
  - `prop(coco, is_a, shibe)`
  - `prop(coco, shibe, true)`
- **Reification** : translating a scenario into object.

## Frames

This can all be brought into a **frame**, collection of attribute-value pairs:

```
[ owned_by = craig
, deliver_to = ming
, model = lemon_laptop_1000 ... ]
```

etc.

## Relations

- **Primitive knowledge** : defined explicitly
- **Derived knowledge** : defined by rules

With an `is_a` attribute, we can do **property inheritance**. Every individual in a class has *n* for some attribute *p*.

No reason not to allow **multiple inheritance**, where an object is a member of multiple classes. There can be conflicts, for example if both classes define a different default for some property (*multiple inheritance problem*).

Associate most general class with an attribute, don't add properties willy-nilly, and axiomatize in the causal direction.

## Complete Knowledge Assumption (CKA)

Any fact not listed in a Knowledge Base is False

The definite clause system is **monotonic**, that is, if we add a clause, it doesn't cause other clauses to be false. (It doesn't invalidate previous conclusions)

Adding the CKA, the system is **non monotonic**; we *can* invalidate a conclusion by adding more clauses.

### Example

```
student(mary).
student(john).
student(ying).
```

- Under the CKA, this means  $\text{student}(X) \iff X = \text{mary} \vee X = \text{john} \vee X = \text{ying}$
- So to prove that  $\neg \text{student}(\text{alan})$ , you need  $\text{alan} \neq \text{mary} \wedge \text{alan} \neq \text{john} \wedge \text{alan} \neq \text{ying}$ 
  - Need unique names assumption

### Clarke Completion

```
mem(X, [X|T]).
mem(X, [H|T]) :- mem(X, T).
```

becomes

```
mem(X, Y) <=> (eT Y == [X | T]) or
               (eH eT T == [H | T] and mem(X, T)).
```

Where  $e$  is  $\exists$

- Completion of every predicate, and equality/inequality axioms.
- If  $p$  in the KB is defined by no clauses, then it completes to  $p \leftrightarrow \text{False}$ . i.e.  $\neg p$ .
- **Negation as Failure** : Interpret negations in clause bodies.  $\sim p$  means  $p$  is False under CKA.

### Bottom up Negation as Failure proof Procedure

```

C = {};
until no more selections possible
  either
    select "h <- b1 or ... or b_n" in KB such that
      bi in C for all i, and h not in C;
    C = C union {h}
  or
    select h such that
      for each "h <- b1 or ... or b_n" in KB
        either
          exists b_i such that ~b_i in C
        or
          exists b_i such that b_i = ~g, and g in C
        end
      C = C union {~h}
    end
  end
end

```

- If this procedure fails, then  $\neg a$  can be concluded.
- Say we have  $a \leftarrow b_1, \dots, a \leftarrow b_n$ , need all of them to fail. It needs to be *finite* however, like  $p \leftarrow p$  is not decidable with bottom up.
  - Halting problem means these things are undetectable in the general case too!
- If trying to NAF a query that has unbound variables, the NAF must be **delayed** until the variable is bound, otherwise it **flounders**.

## Integrity Constraints

- $\text{false} \leftarrow a_1 \wedge \dots \wedge a_k$ 
  - With  $a_i$  atoms, false an atom that's False in all interpretations.
- **Horn Clauses** are either definite clauses or integrity constraints.
- $\neg\alpha$  is a formula which is
  - True in  $I$  if  $\alpha$  is False in  $I$  and
  - False in  $I$  if  $\alpha$  is True in  $I$

Suppose we have KB

- $\text{false} \leftarrow a \wedge b$
- $a \leftarrow c$
- $b \leftarrow c$ 
  - Then  $KB \models \neg c$

Can also be **disjunctive conclusions**:

- $\text{false} \leftarrow a \wedge b$
- $a \leftarrow c$
- $b \leftarrow d$

- Then  $KB \models \neg c \vee \neg d$

### Questions and Answers

- **assumable** : an atom whose negation is acceptably included in the (disjunctive answer)
  - Hand-wavily ‘assumably’ true
- **conflict** : a set of assumables which imply *False* for a KB
- **minimal conflict** : a conflict with no strict subsets that are also conflicts
- **consistency-based diagnosis** : set of assumables that has one element in each conflict
- **minimal diagnosis** : no strict subset is a diagnosis

### Bottom Up Conflict Finding

- **conclusion** : a pair  $\langle a, A \rangle$  with  $a$  an atom, and  $A$  a set of assumables which imply  $a$ .

### Rules and Consistency

- $g$  is ***KB-consistent*** if it’s true in some model of  $KB$ .